

CS 224n Assignment #5

Yoshihiro Kumazawa

July 10, 2020

1. Character-based convolutional encoder for NMT

- (a) Convolutional architectures can operate over variable length input too since convolutional layers slide fixed-sized windows over input unlike linear layers.
- (b) The size of the padding should be 1 so that the padded vector will have size at least 5. Indeed, m_{word} could be 1 if all words in a batch happen to be some characters of length 1 like 'a', in which case we have $\mathbf{x}'_{\text{padded}} \in \mathbb{Z}^3$.
- (c) The highway layer makes it possible to combine local features and global features. In other words, it matches our intuition that we can sometimes understand the meaning of a word by just looking at a little chunk of consecutive characters at a time but it is sometimes better to consider the whole characters in it at once. In order to simplify the network semantics in the beginning of training, I would initialize \mathbf{b}_{gate} to be negative.
- (d) Transformers are easier to parallelize and faster to train.
- (e) See `vocab.py`.
- (f) For the highway network implementation, see `highway.py`. I added a function `question_1f_sanity_check()` in `sanity_check.py` to test the following expected properties.
 - The output size is correct for a given input.
 - $\mathbf{x}_{\text{highway}} = \mathbf{x}_{\text{conv_out}}$ when $\mathbf{x}_{\text{gate}} = 0$, which is checked by making $\mathbf{b}_{\text{gate}} = -\infty$.
 - $\mathbf{x}_{\text{highway}} = \mathbf{x}_{\text{proj}}$ when $\mathbf{x}_{\text{gate}} = 1$, which is checked by making $\mathbf{b}_{\text{gate}} = \infty$.
 - $\mathbf{x}_{\text{highway}} = \mathbf{x}_{\text{conv_out}}$ when the projection layer is the identity function.In addition, I checked if \mathbf{x}_{gate} is initialized to be negative by computing the mean 4 times. (see my answer for 1 (c) above).
- (g) For the convolutional network implementation, see `cnn.py`. I added a function `question_1g_sanity_check()` in `sanity_check.py` to test the following expected properties.
 - The sizes of input channels, output channels, kernels and padding of the convolutional layer are correct.
 - The output size is correct for a given input.
- (h) See `model_embeddings.py`. I do not provide any additional test for it.
- (i)
- (j)

2. Character-based LSTM decoder for NMT

- (a)
- (b)
- (c)
- (d)
- (e)

3. Analyzing NMT Systems

- (a)
- (b)
 - i.
 - ii.
 - iii.
- (c)